

Introduction to Distributed Systems WT 20/21

Assignment 5 – Part I

Submission Deadline: Monday, 08.02.2020, 10:00

- *Submit the solution in PDF via Ilias (only one solution per group).*
- *Respect the submission guidelines (see Ilias).*

1 Two-Phase Locking

[10 points]

Figure 1 shows a serializable history generated by basic (non-strict) Two-Phase Locking for read and write operations originating from four transactions (T_1 , T_2 , T_3 , T_4). In the figure, transaction T_1 is aborted while other transactions T_2 , T_3 , and T_4 are committed.






Figure 1: History H_1

- a) [5 points] Figure 1 shows an incomplete, serializable history of transactions T_1 to T_4 . Add the missing order relationships between the operations in a way, such that serialisability is not violated. Use Figure 2 to show when locks are to be set and released to ensure the serializable history. Also, show the read-/write-accesses in Figure 2. Ensure that all locks are released as early as possible and the operations of each transaction start at the earliest possible time in the space-time diagram, i.e., an operation must not be delayed unless it is blocked due to being unable to acquire locks.
- b) [2 points] In history H_1 shown in Figure 1, some transactions might suffer from cascading abort when transaction T_2 aborts. Name the transactions that are aborted as a result of aborting T_2 . Justify your answer.
- c) [3 points] Now, Strict Two-Phase Locking is used for synchronization. Draw the corresponding space-time diagram (Figure 3). Ensure that the operations of each transaction start at the earliest possible time in the space-time diagram, i.e. an operation must not be delayed unless it is blocked due to being unable to acquire locks. Explain why Strict 2PL prevents cascading aborts?

Figure 2: Basic Two-Phase Locking: Execution of Transactions T_1 to T_4

Figure 3: Strict Two-Phase Locking: Execution of Transactions T_1 to T_4

2 Two-Phase Commit

[11 points]

Figure 4 shows five database servers (nodes) N_1 , N_2 , N_3 , N_4 , and N_5 . Nodes N_1 and N_2 are both located in North America and connected by a wired link. Similarly, nodes N_4 and N_5 reside in Asia and are also connected to each other through a wired link. Node N_3 is located in Europe and communicates with other nodes (N_1 , N_2 , N_4 , and N_5) via two separate satellites.




Figure 4: Connectivity between database servers.

The five nodes perform a distributed transaction initiated by N_3 , i.e., N_3 is the coordinator. To commit the transaction, N_3 starts the 2-Phase Commit Protocol (2PC), as shown in Figure 5. (Note that the space-time diagram in Figure 5 shows the message flow only and no stable states.)

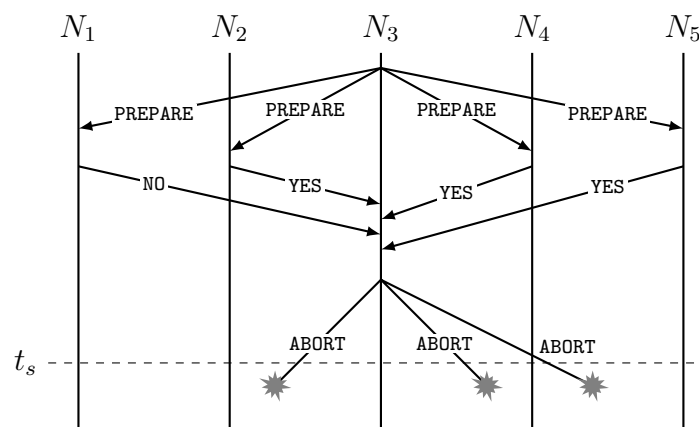


Figure 5: Space-time diagram for the 2-Phase Commit Protocol.

Due to some crash, node N_1 is not prepared to commit the transaction and hence responds with a **NO** vote and aborts the transaction locally. Coordinator node N_3 , hence decides to abort the transaction. At time t_s a heavy thunderstorm above N_3 interrupts the bidirectional link to both the satellites and N_3 is not reachable for several minutes until t_e . Therefore, the **ABORT** messages from N_3 to N_2 , N_4 , and N_5 (nodes which voted **YES** for the transaction) are lost.

- a) [4 points] What happens to N_1 , N_2 , N_3 , N_4 , and N_5 if the *simple termination protocol* is used? Explain the protocol executed by each node until termination of the transaction. Specifically, explain the execution both before and after time t_e . Assume that no node failures occur.
- b) [4 points] What happens to nodes N_1 , N_2 , N_3 , N_4 , and N_5 if the *cooperative termination protocol* is used? Sketch the execution of the protocol until termination of the transaction by completing the space-time diagram from Fig. 5, showing all message exchanges. Specifically, explain the execution both before and after time t_e . Also, indicate for nodes N_1 , N_2 , N_4 , and N_5 whether they remain blocked or not until the connectivity to node N_3 is re-established (i.e., time t_e is reached). Again, assume that no node failures occur.
- Note: we consider a node being “blocked” as long as it is unable to unilaterally abort a running transaction, i.e. it has sent a YES vote but is still waiting for a termination decision (COMMIT or ABORT).*
- c) [3 points] Now assume that N_4 suffers a crash failure immediately after sending its YES message to N_3 and recovers after t_e . Modify your diagram from part (b) accordingly, under the assumption of the failure of N_4 as mentioned above.

3 Data Replication

[13 points]

Given are four nodes A , B , C , D as shown in Figure 6. An object X is replicated to these nodes using the “weighted voting” replication protocol. Each node n is available with a certain probability p_n , also shown in the figure, and maintains a replica X_n of the logical object X . Assume that no communication failures occur (i.e., network partitioning is not possible).




Figure 6: Configuration for Weighted Voting

- a) [5 points] Assume that read and write thresholds of $q_r[X] = 4$ and $q_w[X] = 5$ are given. Find a valid assignment of values for $q[X_A]$, $q[X_B]$, $q[X_C]$, and $q[X_D]$ that
1. minimizes the number of nodes that have to be locked for writing X
 2. maximizes the availability of X for writing

Both conditions need to be fulfilled (condition 1 with a higher priority). Also, give all possible read and write quorums for your solution.

Note: each node must have at least one vote, i.e. $q[X_n] \geq 1$. Consider all quorums, not only minimal ones.

- b) [3 points] Now, find a valid assignment of the values for $q[X_A]$, $q[X_B]$, $q[X_C]$, and $q[X_D]$ yielding a higher write-availability than the solution of task a), when ignoring

condition 1. Assume the same thresholds as in the previous task and every replica needs at least one vote. State the provided write-availability. Justify your answer.

- c) [3 points] Now, assume that $q[X_A] = 3, q[X_B] = 3, q[X_C] = 1, q[X_D] = 2$. Find a write threshold $q_w[X]$ that enables exactly 6 different write quorums. Give the threshold value $q_w[X]$ and the resulting write quorums.

Note: Consider all quorums, not only minimal quorums.

- d) [2 points] *Majority Consensus* is a special case of Weighted Voting replication management protocol. Now, assume that the *Majority Consensus* is used to replicate an object Y in a system with three nodes K, L, M where each node has exactly 1 vote. As before, each node n is available with a certain probability p_n . Let $p_r(Y)$ be the probability that the logical object Y is available for *reading*. What is $p_r(Y)$ in terms of p_K, p_L , and p_M ? Justify your answer.

